

prac_linear

April 29, 2025

```
[1]: import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from tensorflow.keras.datasets import boston_housing
import tensorflow as tf
from tensorflow.keras import layers, models
import matplotlib.pyplot as plt
```

```
2025-04-29 07:14:28.590578: I external/local_xla/xla/tsl/cuda/cudart_stub.cc:32]
Could not find cuda drivers on your machine, GPU will not be used.
2025-04-29 07:14:28.594512: I external/local_xla/xla/tsl/cuda/cudart_stub.cc:32]
Could not find cuda drivers on your machine, GPU will not be used.
2025-04-29 07:14:28.607916: E
external/local_xla/xla/stream_executor/cuda/cuda_fft.cc:477] Unable to register
cuFFT factory: Attempting to register factory for plugin cuFFT when one has
already been registered
WARNING: All log messages before absl::InitializeLog() is called are written to
STDERR
E0000 00:00:1745925268.630763    10212 cuda_dnn.cc:8310] Unable to register cuDNN
factory: Attempting to register factory for plugin cuDNN when one has already
been registered
E0000 00:00:1745925268.637633    10212 cuda_blas.cc:1418] Unable to register
cuBLAS factory: Attempting to register factory for plugin cuBLAS when one has
already been registered
2025-04-29 07:14:28.660683: I tensorflow/core/platform/cpu_feature_guard.cc:210]
This TensorFlow binary is optimized to use available CPU instructions in
performance-critical operations.
To enable the following instructions: AVX2 FMA, in other operations, rebuild
TensorFlow with the appropriate compiler flags.
```

```
[2]: (X_train, y_train), (X_test, y_test) = boston_housing.load_data()
```

```
[3]: X = np.vstack((X_train, X_test))
y = np.hstack((y_train, y_test))
```

```
[4]: feature_names = [
      'CRIM', 'ZN', 'INDUS', 'CHAS', 'NOX', 'RM',
      'AGE', 'DIS', 'RAD', 'TAX', 'PTRATIO', 'B', 'LSTAT'
    ]
```

```
[5]: df = pd.DataFrame(X, columns=feature_names)
      df['Price'] = y
```

```
[6]: print(df.head())
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	\
0	1.23247	0.0	8.14	0.0	0.538	6.142	91.7	3.9769	4.0	307.0	
1	0.02177	82.5	2.03	0.0	0.415	7.610	15.7	6.2700	2.0	348.0	
2	4.89822	0.0	18.10	0.0	0.631	4.970	100.0	1.3325	24.0	666.0	
3	0.03961	0.0	5.19	0.0	0.515	6.037	34.5	5.9853	5.0	224.0	
4	3.69311	0.0	18.10	0.0	0.713	6.376	88.4	2.5671	24.0	666.0	

	PTRATIO	B	LSTAT	Price
0	21.0	396.90	18.72	15.2
1	14.7	395.38	3.11	42.3
2	20.2	375.52	3.26	50.0
3	20.2	396.90	8.01	21.1
4	20.2	391.43	14.65	17.7

```
[7]: # Features and target variable
      X = df.drop('Price', axis=1).values
      y = df['Price'].values
```

```
[8]: # Standardize the data (important for neural networks)
      scaler = StandardScaler()
      X_train = scaler.fit_transform(X_train)
      X_test = scaler.transform(X_test)
```

```
[9]: # Build the Deep Neural Network model for Linear Regression
      model = models.Sequential([
          layers.Dense(64, activation='relu', input_dim=X_train.shape[1]),
          layers.Dense(64, activation='relu'),
          layers.Dense(32, activation='relu'),
          layers.Dense(1) # Output layer with 1 neuron for regression output (no
                           ↪ activation function)
      ])
```

```
/home/anil/.local/lib/python3.10/site-
packages/keras/src/layers/core/dense.py:87: UserWarning: Do not pass an
`input_shape`/`input_dim` argument to a layer. When using Sequential models,
prefer using an `Input(shape)` object as the first layer in the model instead.
  super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

W0000 00:00:1745925271.798467 10212 gpu_device.cc:2344] Cannot dlopen some GPU libraries. Please make sure the missing libraries mentioned above are installed properly if you would like to use GPU. Follow the guide at <https://www.tensorflow.org/install/gpu> for how to download and setup the required libraries for your platform.
Skipping registering GPU devices...

```
[10]: # Compile the model
model.compile(optimizer='adam', loss='mean_squared_error')
```

```
[11]: # Train the model
history = model.fit(X_train, y_train, epochs=10, batch_size=32,
    ↪ validation_data=(X_test, y_test), verbose=1)
```

```
Epoch 1/10
13/13          1s 14ms/step -
loss: 603.2776 - val_loss: 604.9301
Epoch 2/10
13/13          0s 4ms/step - loss:
565.3013 - val_loss: 568.8455
Epoch 3/10
13/13          0s 4ms/step - loss:
505.8569 - val_loss: 499.5744
Epoch 4/10
13/13          0s 3ms/step - loss:
446.8285 - val_loss: 370.1459
Epoch 5/10
13/13          0s 3ms/step - loss:
306.5861 - val_loss: 182.2021
Epoch 6/10
13/13          0s 4ms/step - loss:
130.7779 - val_loss: 62.5338
Epoch 7/10
13/13          0s 4ms/step - loss:
56.3639 - val_loss: 48.4146
Epoch 8/10
13/13          0s 4ms/step - loss:
39.7146 - val_loss: 36.6643
Epoch 9/10
13/13          0s 4ms/step - loss:
29.9657 - val_loss: 30.5572
Epoch 10/10
13/13          0s 3ms/step - loss:
24.9629 - val_loss: 27.6493
```

```
[12]: # Evaluate the model on the test data
test_loss = model.evaluate(X_test, y_test)
```


```
print(f"Test Loss (Mean Squared Error): {test_loss}")
```

```
4/4          0s 2ms/step - loss:
25.3579
Test Loss (Mean Squared Error): 27.649261474609375
```


```
[13]: # Make predictions
predictions = model.predict(X_test)
```

```
4/4          0s 14ms/step
```

```
[14]: # Compare the predictions to the actual values
plt.figure(figsize=(8, 6))
plt.scatter(y_test, predictions)
plt.plot([min(y_test), max(y_test)], [min(y_test), max(y_test)], color='red')
    ↪ # 45-degree line
plt.xlabel('True Values')
plt.ylabel('Predictions')
plt.title('True vs Predicted Housing Prices')
plt.show()
```



```
[15]: # Optionally, plot the loss curve
plt.plot(history.history['loss'], label='Train Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.title('Loss Curve')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.show()
```




```
[16]: import seaborn as sns
print("\nCorrelation matrix:")
correlation_matrix = df.corr()
print(correlation_matrix)

plt.figure(figsize=(12, 8))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt='.2f')
plt.title('Correlation Matrix of Features')
plt.show()
```

Correlation matrix:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	\
CRIM	1.000000	-0.200469	0.406583	-0.055892	0.420972	-0.219247	0.352734	
ZN	-0.200469	1.000000	-0.533828	-0.042697	-0.516604	0.311991	-0.569537	
INDUS	0.406583	-0.533828	1.000000	0.062938	0.763651	-0.391676	0.644779	
CHAS	-0.055892	-0.042697	0.062938	1.000000	0.091203	0.091251	0.086518	
NOX	0.420972	-0.516604	0.763651	0.091203	1.000000	-0.302188	0.731470	
RM	-0.219247	0.311991	-0.391676	0.091251	-0.302188	1.000000	-0.240265	
AGE	0.352734	-0.569537	0.644779	0.086518	0.731470	-0.240265	1.000000	
DIS	-0.379670	0.664408	-0.708027	-0.099176	-0.769230	0.205246	-0.747881	
RAD	0.625505	-0.311948	0.595129	-0.007368	0.611441	-0.209847	0.456022	
TAX	0.582764	-0.314563	0.720760	-0.035587	0.668023	-0.292048	0.506456	
PTRATIO	0.289946	-0.391679	0.383248	-0.121515	0.188933	-0.355501	0.261515	
B	-0.385064	0.175520	-0.356977	0.048788	-0.380051	0.128069	-0.273534	
LSTAT	0.455621	-0.412995	0.603800	-0.053929	0.590879	-0.613808	0.602339	
Price	-0.388305	0.360445	-0.483725	0.175260	-0.427321	0.695360	-0.376955	

	DIS	RAD	TAX	PTRATIO	B	LSTAT	Price
CRIM	-0.379670	0.625505	0.582764	0.289946	-0.385064	0.455621	-0.388305
ZN	0.664408	-0.311948	-0.314563	-0.391679	0.175520	-0.412995	0.360445
INDUS	-0.708027	0.595129	0.720760	0.383248	-0.356977	0.603800	-0.483725
CHAS	-0.099176	-0.007368	-0.035587	-0.121515	0.048788	-0.053929	0.175260
NOX	-0.769230	0.611441	0.668023	0.188933	-0.380051	0.590879	-0.427321
RM	0.205246	-0.209847	-0.292048	-0.355501	0.128069	-0.613808	0.695360
AGE	-0.747881	0.456022	0.506456	0.261515	-0.273534	0.602339	-0.376955
DIS	1.000000	-0.494588	-0.534432	-0.232471	0.291512	-0.496996	0.249929
RAD	-0.494588	1.000000	0.910228	0.464741	-0.444413	0.488676	-0.381626
TAX	-0.534432	0.910228	1.000000	0.460853	-0.441808	0.543993	-0.468536
PTRATIO	-0.232471	0.464741	0.460853	1.000000	-0.177383	0.374044	-0.507787
B	0.291512	-0.444413	-0.441808	-0.177383	1.000000	-0.366087	0.333461
LSTAT	-0.496996	0.488676	0.543993	0.374044	-0.366087	1.000000	-0.737663
Price	0.249929	-0.381626	-0.468536	-0.507787	0.333461	-0.737663	1.000000



[17]: *# Visualizing the distribution of the target variable (Price)*

```
plt.figure(figsize=(8, 6))
sns.histplot(df['Price'], kde=True, bins=30)
plt.title('Distribution of Housing Prices')
plt.xlabel('Price')
plt.ylabel('Frequency')
plt.show()
```

ValueError Traceback (most recent call last)

Cell In[17], line 3

```
1 # Visualizing the distribution of the target variable (Price)
2 plt.figure(figsize=(8, 6))
----> 3 sns.histplot(df['Price'], kde=True, bins=30)
4 plt.title('Distribution of Housing Prices')
5 plt.xlabel('Price')
```

File ~/local/lib/python3.10/site-packages/seaborn/distributions.py:1416, in
 ↪ histplot(data, x, y, hue, weights, stat, bins, binwidth, binrange, discrete,
 ↪ cumulative, common_bins, common_norm, multiple, element, fill, shrink, kde,
 ↪ kde_kws, line_kws, thresh, pthresh, pmax, cbar, cbar_ax, cbar_kws, palette,
 ↪ hue_order, hue_norm, color, log_scale, legend, ax, **kwargs)

```

1405 estimate_kws = dict(
1406     stat=stat,
1407     bins=bins,
1408     (...)
1411     cumulative=cumulative,
1412 )
1414 if p.univariate:
-> 1416     p.plot_univariate_histogram(
1417         multiple=multiple,
1418         element=element,
1419         fill=fill,
1420         shrink=shrink,
1421         common_norm=common_norm,
1422         common_bins=common_bins,
1423         kde=kde,
1424         kde_kws=kde_kws,
1425         color=color,
1426         legend=legend,
1427         estimate_kws=estimate_kws,
1428         line_kws=line_kws,
1429         **kwargs,
1430     )
1432 else:
1434     p.plot_bivariate_histogram(
1435         common_bins=common_bins,
1436         common_norm=common_norm,
1437         (...)
1446         **kwargs,
1447     )

```

File ~/.local/lib/python3.10/site-packages/seaborn/distributions.py:651, in `DistributionPlotter.plot_univariate_histogram`(self, multiple, element, fill, common_norm, common_bins, shrink, kde, kde_kws, color, legend, line_kws, estimate_kws, **plot_kws)

```

648     sticky_x, sticky_y = (0, np.inf), None
650     line_kws["color"] = to_rgba(sub_color, 1)
--> 651     line, = ax.plot(
652         *line_args, **line_kws,
653     )
655     if sticky_x is not None:
656         line.sticky_edges.x[:] = sticky_x

```

File /usr/lib/python3/dist-packages/matplotlib/axes/_axes.py:1632, in `Axes.plot`(self, scalex, scaley, data, *args, **kwargs)

```

1390 """
1391 Plot y versus x as lines and/or markers.
1392 (...)

```



```

1629 (``'green'``) or hex strings (``'#008000'``).
1630 """
1631 kwargs = cbook.normalize_kwargs(kwargs, mlines.Line2D)
-> 1632 lines = [*self._get_lines(*args, data=data, **kwargs)]
1633 for line in lines:
1634     self.add_line(line)

```

```

File /usr/lib/python3/dist-packages/matplotlib/axes/_base.py:312, in
-> _process_plot_var_args.__call__(self, data, *args, **kwargs)
    310     this += args[0],
    311     args = args[1:]
--> 312 yield from self._plot_args(this, kwargs)

```

```

File /usr/lib/python3/dist-packages/matplotlib/axes/_base.py:487, in
-> _process_plot_var_args._plot_args(self, tup, kwargs, return_kwargs)
    484     kw[prop_name] = val
    486 if len(xy) == 2:
--> 487     x = _check_1d(xy[0])
    488     y = _check_1d(xy[1])
    489 else:

```

```

File /usr/lib/python3/dist-packages/matplotlib/cbook/__init__.py:1327, in
-> _check_1d(x)
    1321 with warnings.catch_warnings(record=True) as w:
    1322     warnings.filterwarnings(
    1323         "always",
    1324         category=Warning,
    1325         message='Support for multi-dimensional indexing')
-> 1327     ndim = x[:, None].ndim
    1328     # we have definitely hit a pandas index or series object
    1329     # cast to a numpy array.
    1330     if len(w) > 0:

```

```

File ~/.local/lib/python3.10/site-packages/pandas/core/indexes/base.py:5419, in
-> Index.__getitem__(self, key)
    5417 # Because we ruled out integer above, we always get an arraylike here
    5418 if result.ndim > 1:
-> 5419     disallow_ndim_indexing(result)
    5421 # NB: Using _constructor._simple_new would break if MultiIndex
    5422 # didn't override __getitem__
    5423 return self._constructor._simple_new(result, name=self._name)

```

```

File ~/.local/lib/python3.10/site-packages/pandas/core/indexers/utils.py:341, in
-> disallow_ndim_indexing(result)
    333 """
    334 Helper function to disallow multi-dimensional indexing on 1D Series/
-> Index.
    335


```

```

(...)
338 in GH#30588.
339 """
340 if np.ndim(result) > 1:
--> 341     raise ValueError(
342         "Multi-dimensional indexing (e.g. `obj[:, None]`) is no longer
343         "supported. Convert to a numpy array before indexing instead."
344     )

```


ValueError: Multi-dimensional indexing (e.g. `obj[:, None]`) is no longer supported. Convert to a numpy array before indexing instead.



```

[18]: # Visualizing the relationship between the most correlated feature and Price
plt.figure(figsize=(8, 6))
sns.scatterplot(x=df['RM'], y=df['Price'])
plt.title('Price vs Average Number of Rooms (RM)')
plt.xlabel('Average Number of Rooms (RM)')
plt.ylabel('Price')
plt.show()

```



[]: